

Shipt – Interview Challenge – Data Analyst

Required Questions

```
In [321]: # Importing the libraries
import pandas as pd
import numpy as np

import warnings
from pandas.core.common import SettingWithCopyWarning
warnings.simplefilter(action="ignore", category=SettingWithCopyWarning)
```

```
In [322]: # Reading the InterviewData_Cost.csv file
InterviewData_Cost = pd.read_csv('/Users/abhilashasinha/Downloads/interv
iew_challenge_data_analyst/InterviewData_Cost.csv')
```

```
In [323]: # Data in InterviewData_Cost
InterviewData_Cost
```

Out[323]:

	date	source_id	cost
0	10/17/14	PA0577	7168.0
1	8/17/14	PA0354	7615.0
2	1/7/14	PA0607	4054.0
3	8/25/14	PA0745	9317.0
4	11/30/14	PA0923	5586.0
...
9995	1/10/14	PA0830	6202.0
9996	1/31/14	PA0467	7057.0
9997	10/14/14	PA0277	9112.0
9998	10/1/14	PA0293	1053.0
9999	1/29/14	PA0470	4186.0

10000 rows × 3 columns

```
In [324]: # Reading the InterviewData_Rev.csv file
InterviewData_Rev = pd.read_csv('/Users/abhilashasinha/Downloads/intervi
ew_challenge_data_analyst/InterviewData_Rev.csv')
```

```
In [325]: # Data in InterviewData_Rev  
InterviewData_Rev
```

Out[325]:

	date	source_id	revenue
0	8/1/14	PA0368	5717.0
1	1/31/14	PA0277	1380.0
2	6/9/14	PA0745	7535.0
3	9/1/14	PA0751	2868.0
4	3/12/14	PA0859	10757.0
...
9995	9/24/14	PA0174	3827.0
9996	10/26/14	PA0318	533.0
9997	8/3/14	PA0923	4963.0
9998	1/4/14	PA0354	3070.0
9999	6/7/14	PA0354	6716.0

10000 rows × 3 columns

Q1. Join these two data sets by “date” and “source_id”, returning all rows from both regardless of whether there is a match between the two data sets

Ans: Since we need all rows from both InterviewData_Cost and InterviewData_Rev regardless of whether there is a match between the two data sets, I have joined them using outer join on 'date' and 'source_id'.

```
In [326]: # Join the two tables using outer join
data_all = pd.merge(InterviewData_Cost, InterviewData_Rev, on=['date', 'source_id'], how='outer')
data_all
```

Out[326]:

	date	source_id	cost	revenue
0	10/17/14	PA0577	7168.0	8417.0
1	8/17/14	PA0354	7615.0	4200.0
2	1/7/14	PA0607	4054.0	7935.0
3	8/25/14	PA0745	9317.0	5536.0
4	11/30/14	PA0923	5586.0	NaN
...
14613	12/29/14	PA0958	NaN	7406.0
14614	3/19/14	PA0732	NaN	6584.0
14615	9/24/14	PA0174	NaN	3827.0
14616	10/26/14	PA0318	NaN	533.0
14617	1/4/14	PA0354	NaN	3070.0

14618 rows × 4 columns

```
In [327]: data_all['source_id'].unique()
```

```
Out[327]: array(['PA0577', 'PA0354', 'PA0607', 'PA0745', 'PA0923', 'PA0808',
                'PA0952', 'PA0411', 'PA0526', 'PA0168', 'PA0277', 'PA0057',
                'PA0482', 'PA0368', 'PA0552', 'PA0696', 'PA0258', 'PA0338',
                'PA0619', 'PA0270', 'PA0474', 'PA0958', 'PA0308', 'PA0830',
                'PA0752', 'PA0859', 'PA0174', 'PA0792', 'PA0732', 'PA0318',
                'PA0293', 'PA0198', 'PA0202', 'PA0843', 'PA0873', 'PA0973',
                'PA0394', 'PA0751', 'PA0169', 'PA0352', 'PA0126', 'PA0672',
                'PA0900', 'PA0527', 'PA0534', 'PA0470', 'PA0543', 'PA0467',
                'PA0678', 'PA0659'], dtype=object)
```

Q2. Join these two data sets by “date” and “source_id”, returning only the rows from the “Cost” file that have no corresponding date in the “Revenue” file.

Ans: Here, I have used left join to join the two data sets. As we need all the rows from InterviewData_Cost, it is placed on the left side of the query.

```
In [328]: # Join the two tables using left join
pd.merge(InterviewData_Cost, InterviewData_Rev, on=[ 'date', 'source_id' ], how='left')
```

Out[328]:

	date	source_id	cost	revenue
0	10/17/14	PA0577	7168.0	8417.0
1	8/17/14	PA0354	7615.0	4200.0
2	1/7/14	PA0607	4054.0	7935.0
3	8/25/14	PA0745	9317.0	5536.0
4	11/30/14	PA0923	5586.0	NaN
...
9995	1/10/14	PA0830	6202.0	NaN
9996	1/31/14	PA0467	7057.0	NaN
9997	10/14/14	PA0277	9112.0	8853.0
9998	10/1/14	PA0293	1053.0	NaN
9999	1/29/14	PA0470	4186.0	2146.0

10000 rows × 4 columns

Q3. Using your result from #1:

a. What are the Top 4 sources (“source_id” values) in terms of total revenue generation across this data set?

b. How would you visualize the monthly revenue for those Top 4 sources?

a. What are the Top 4 sources (“source_id” values) in terms of total revenue generation across this data set?

```
In [329]: # data_all_total contains the aggregate of revenue for each source_id  
data_all_total = data_all.groupby(data_all['source_id']).agg({'revenue':  
    'sum'}).reset_index()  
data_all_total
```

Out[329]:

	source_id	revenue
0	PA0057	1032845.0
1	PA0126	1245754.0
2	PA0168	1132637.0
3	PA0169	1223275.0
4	PA0174	1203643.0
5	PA0198	1152959.0
6	PA0202	1246337.0
7	PA0258	1185512.0
8	PA0270	1237148.0
9	PA0277	1201316.0
10	PA0293	1195860.0
11	PA0308	1338615.0
12	PA0318	1242835.0
13	PA0338	1142116.0
14	PA0352	1309685.0
15	PA0354	1102736.0
16	PA0368	1147076.0
17	PA0394	1138108.0
18	PA0411	1130753.0
19	PA0467	1103085.0
20	PA0470	1148255.0
21	PA0474	1169431.0
22	PA0482	1217803.0
23	PA0526	1257963.0
24	PA0527	1385747.0
25	PA0534	1254734.0
26	PA0543	1266861.0
27	PA0552	1283190.0
28	PA0577	1244647.0
29	PA0607	1236800.0
30	PA0619	1110368.0
31	PA0659	1228329.0
32	PA0672	1138589.0
33	PA0678	955864.0
34	PA0696	1279198.0

	source_id	revenue
35	PA0732	1205794.0
36	PA0745	1053657.0
37	PA0751	1240378.0
38	PA0752	1185995.0
39	PA0792	1159598.0
40	PA0808	1026800.0
41	PA0830	1237014.0
42	PA0843	1204455.0
43	PA0859	1173593.0
44	PA0873	1211288.0
45	PA0900	1243498.0
46	PA0923	1236477.0
47	PA0952	1112659.0
48	PA0958	1271398.0
49	PA0973	1165452.0

```
In [330]: #Sorting the data obtained from previous query in descending order of to  
tal revenue  
data_all_total.sort_values(by=['revenue'], inplace=True, ascending=False  
)  
data_all_total
```


Out[330]:

	source_id	revenue
24	PA0527	1385747.0
11	PA0308	1338615.0
14	PA0352	1309685.0
27	PA0552	1283190.0
34	PA0696	1279198.0
48	PA0958	1271398.0
26	PA0543	1266861.0
23	PA0526	1257963.0
25	PA0534	1254734.0
6	PA0202	1246337.0
1	PA0126	1245754.0
28	PA0577	1244647.0
45	PA0900	1243498.0
12	PA0318	1242835.0
37	PA0751	1240378.0
8	PA0270	1237148.0
41	PA0830	1237014.0
29	PA0607	1236800.0
46	PA0923	1236477.0
31	PA0659	1228329.0
3	PA0169	1223275.0
22	PA0482	1217803.0
44	PA0873	1211288.0
35	PA0732	1205794.0
42	PA0843	1204455.0
4	PA0174	1203643.0
9	PA0277	1201316.0
10	PA0293	1195860.0
38	PA0752	1185995.0
7	PA0258	1185512.0
43	PA0859	1173593.0
21	PA0474	1169431.0
49	PA0973	1165452.0
39	PA0792	1159598.0
5	PA0198	1152959.0

	source_id	revenue
20	PA0470	1148255.0
16	PA0368	1147076.0
13	PA0338	1142116.0
32	PA0672	1138589.0
17	PA0394	1138108.0
2	PA0168	1132637.0
18	PA0411	1130753.0
47	PA0952	1112659.0
30	PA0619	1110368.0
19	PA0467	1103085.0
15	PA0354	1102736.0
36	PA0745	1053657.0
0	PA0057	1032845.0
40	PA0808	1026800.0
33	PA0678	955864.0

```
In [331]: #Top 4 source_ids with highest revenue
data_all_total_top4 = data_all_total.head(4)
data_all_total_top4
```

Out[331]:

	source_id	revenue
24	PA0527	1385747.0
11	PA0308	1338615.0
14	PA0352	1309685.0
27	PA0552	1283190.0

```
In [332]: # Creating a list with only the top 4 source_ids
data_source_id4 = data_all_total_top4['source_id'].tolist()
```

Top 4 sources (“source_id” values) in terms of total revenue generation across this data set

```
In [333]: # Data in the list data_source_id4 showing top 4 “source_id” values in t
           erms of total revenue generation across data set
data_source_id4
```

Out[333]: ['PA0527', 'PA0308', 'PA0352', 'PA0552']

b. How would you visualize the monthly revenue for those Top 4 sources?

In order to visualize the monthly revenue, I have created a dataframe having the top 4 source_ids and the aggregated revenue generated in each month.

```
In [341]: # Selecting records which belongs to the top 4 source_ids having highest revenue
data_all_filtered = data_all[data_all['source_id'].isin(data_source_id4)]
```

```
In [342]: data_all_filtered
```

```
Out[342]:
```

	date	source_id	cost	revenue
16	3/29/14	PA0552	7894.0	NaN
22	8/21/14	PA0552	4653.0	NaN
27	4/28/14	PA0308	3958.0	10863.0
65	12/20/14	PA0552	2092.0	NaN
76	8/13/14	PA0308	777.0	3863.0
...
14543	7/26/14	PA0308	NaN	1262.0
14545	7/2/14	PA0308	NaN	3582.0
14558	10/14/14	PA0308	NaN	10001.0
14580	7/5/14	PA0308	NaN	11646.0
14594	10/22/14	PA0352	NaN	10070.0

1187 rows × 4 columns

```
In [343]: # Unique source_ids
data_all_filtered.source_id.unique()
```

```
Out[343]: array(['PA0552', 'PA0308', 'PA0352', 'PA0527'], dtype=object)
```

```
In [344]: # Checking the data types
data_all_filtered.dtypes
```

```
Out[344]: date          object
source_id      object
cost           float64
revenue        float64
dtype: object
```

```
In [345]: # Converting date from object to date format
data_all_filtered['date'] = pd.to_datetime(data_all_filtered['date'], format = '%m/%d/%y')
```

According to the question, we need the monthly revenue, so I have selected the sum of revenue for each source_id in each month in year 2014.

```
In [346]: # data_all_date contains the monthly aggregate of revenue for each source_id
data_all_date = data_all_filtered.groupby([data_all_filtered['date'].dt.month, data_all_filtered['source_id']]).agg({'revenue': 'sum'})
```

```
In [347]: data_all_date
```

Out[347]:

revenue		
date	source_id	
1	PA0308	74573.0
	PA0352	103252.0
	PA0527	111528.0
	PA0552	95687.0
2	PA0308	108134.0
	PA0352	83001.0
	PA0527	105195.0
	PA0552	99742.0
3	PA0308	142817.0
	PA0352	100480.0
	PA0527	135439.0
	PA0552	90240.0
4	PA0308	164805.0
	PA0352	93279.0
	PA0527	112385.0
	PA0552	111432.0
5	PA0308	122349.0
	PA0352	148362.0
	PA0527	114339.0
	PA0552	129668.0
6	PA0308	76398.0
	PA0352	139655.0
	PA0527	112022.0
	PA0552	103569.0
7	PA0308	104544.0
	PA0352	118855.0
	PA0527	128733.0
	PA0552	117613.0
8	PA0308	99909.0
	PA0352	115949.0
	PA0527	91288.0
	PA0552	78776.0
9	PA0308	99915.0
	PA0352	117853.0

		revenue
date	source_id	
	PA0527	140484.0
	PA0552	118337.0
	PA0308	113668.0
	PA0352	100393.0
10	PA0527	135752.0
	PA0552	102767.0
	PA0308	107097.0
	PA0352	78840.0
11	PA0527	66297.0
	PA0552	94062.0
	PA0308	124406.0
	PA0352	109766.0
12	PA0527	132285.0
	PA0552	141297.0

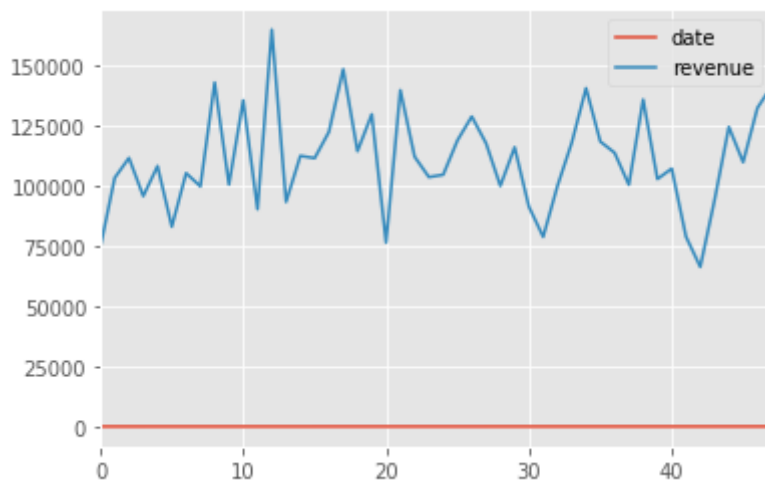
Visualizing the monthly revenue for those Top 4 sources

```
In [348]: data_all_vis = data_all_date.reset_index()
```

```
In [349]: # Plot to visualize the data
import matplotlib
import matplotlib.pyplot as plt
import seaborn as sns

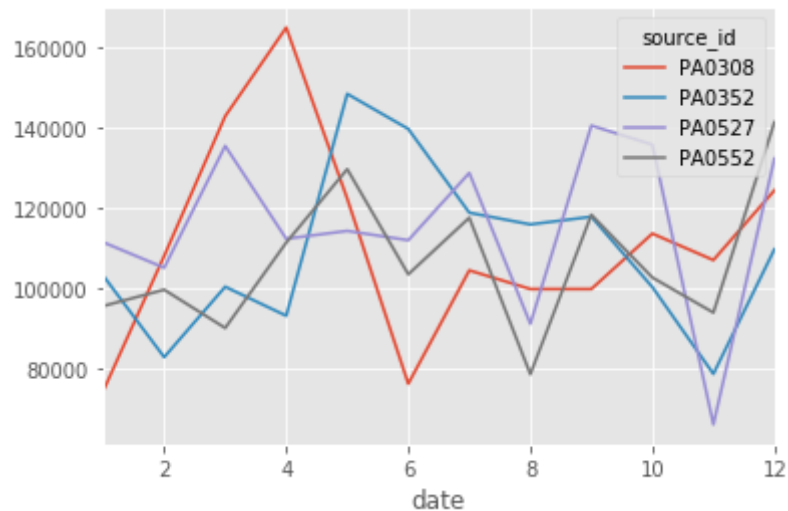
data_all_vis.plot()
```

```
Out[349]: <matplotlib.axes._subplots.AxesSubplot at 0x1c282677d0>
```



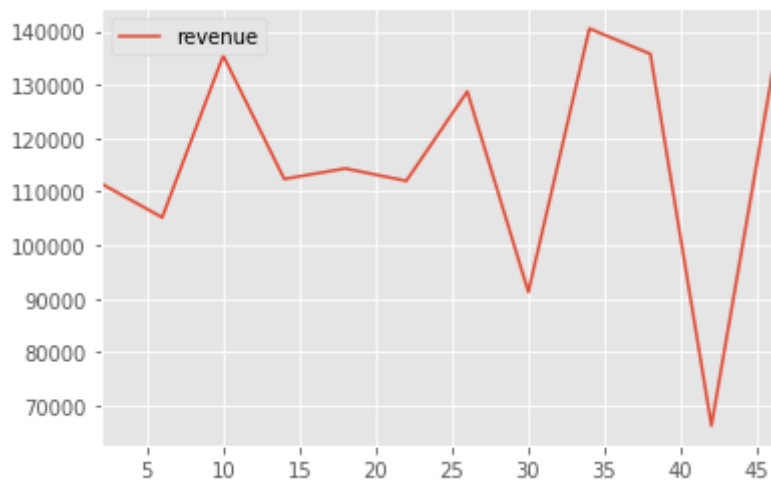
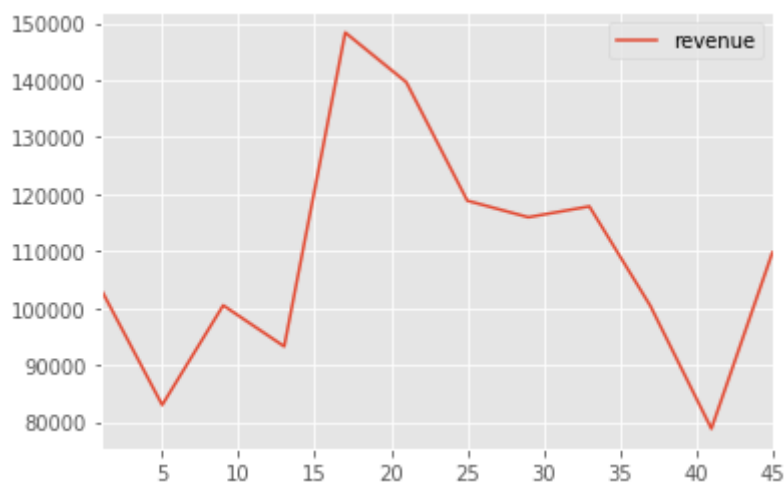
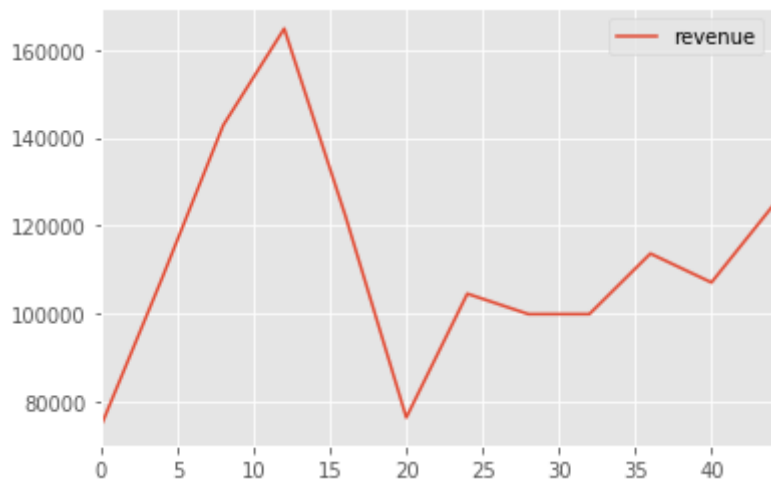
```
In [350]: # Monthly revenue for Top 4 sources plotted together
pv = pd.pivot_table(data_all_vis, index=data_all_vis.date, columns=data_all_vis.source_id,
                    values='revenue')
pv.plot()
```

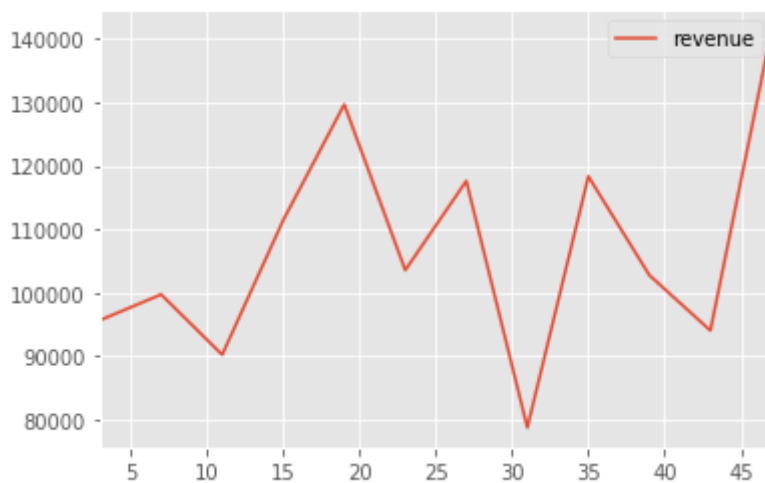
Out[350]: <matplotlib.axes._subplots.AxesSubplot at 0x11a679a10>




```
In [351]: # Single plot for each source_id to understand the revenue pattern of ea  
ch source  
data_all_vis.groupby('source_id').plot(y='revenue')
```

```
Out[351]: source_id
PA0308    AxesSubplot(0.125,0.125;0.775x0.755)
PA0352    AxesSubplot(0.125,0.125;0.775x0.755)
PA0527    AxesSubplot(0.125,0.125;0.775x0.755)
PA0552    AxesSubplot(0.125,0.125;0.775x0.755)
dtype: object
```





Questions 4 and 5 deal with “InterviewData_Activity.csv”.

4. Assuming you’ve read the data into an R object called `activity_data`, run the following code to build a basic logistic regression model:

```
In [352]: # Reading the InterviewData_Activity.csv file
activity_data = pd.read_csv('/Users/abhilashasinha/Downloads/interview_challenge_data_analyst/InterviewData_Activity.csv')
```

```
In [353]: # Data in activity_data
activity_data
```

Out[353]:

	userid	date	age	gender	metropolitan_area	device_type	active
0	4e3a9ea937b3a	8/4/15	30	F	Detroit	Tablet	1
1	4e3dd5154a08c	8/6/15	43	F	Charlotte	Desktop	1
2	4e3df1ecd131a	8/6/15	41	F	Tampa	Mobile	1
3	4e4e77461b1e3	8/19/15	56	F	Nashville	Desktop	1
4	4e4eb59b6de55	8/19/15	33	F	Detroit	Mobile	1
...
5415	4e9cce2b66d44	10/17/15	21	M	Houston	Tablet	0
5416	4f2ab00039f50	2/2/16	42	F	Birmingham	Mobile	0
5417	4e642e7208f6b	9/4/15	40	F	Houston	Mobile	0
5418	4f40149878765	2/18/16	40	F	Houston	Mobile	0
5419	4ea84c57de3bb	10/26/15	30	F	Tampa	Mobile	0

5420 rows × 7 columns

```
In [354]: # Importing the library for statistical model
import statsmodels.api as sm
```

```
In [355]: dummy_genders = pd.get_dummies(activity_data['gender'], prefix = 'gender')
dummy_metro = pd.get_dummies(activity_data['metropolitan_area'], prefix = 'metro_area')
dummy_device = pd.get_dummies(activity_data['device_type'], prefix = 'device')
```

```
In [356]: activity_data
```

```
Out[356]:
```

	userid	date	age	gender	metropolitan_area	device_type	active
0	4e3a9ea937b3a	8/4/15	30	F	Detroit	Tablet	1
1	4e3dd5154a08c	8/6/15	43	F	Charlotte	Desktop	1
2	4e3df1ecd131a	8/6/15	41	F	Tampa	Mobile	1
3	4e4e77461b1e3	8/19/15	56	F	Nashville	Desktop	1
4	4e4eb59b6de55	8/19/15	33	F	Detroit	Mobile	1
...
5415	4e9cce2b66d44	10/17/15	21	M	Houston	Tablet	0
5416	4f2ab00039f50	2/2/16	42	F	Birmingham	Mobile	0
5417	4e642e7208f6b	9/4/15	40	F	Houston	Mobile	0
5418	4f40149878765	2/18/16	40	F	Houston	Mobile	0
5419	4ea84c57de3bb	10/26/15	30	F	Tampa	Mobile	0

5420 rows × 7 columns

```
In [357]: cols_to_keep = ['active', 'age']
```

```
In [358]: activity_data = activity_data[cols_to_keep].join(dummy_genders.ix[:, 'gender_M':])
activity_data = activity_data.join(dummy_metro.ix[:, 'metro_area_Birmingham':])
activity_data = activity_data.join(dummy_device.ix[:, 'device_Mobile':])
activity_data = sm.add_constant(activity_data, prepend=False)
```

```
In [359]: explanatory_cols = activity_data.columns[1:]
full_logit_model = sm.GLM(activity_data['active'], activity_data[explanatory_cols], family=sm.families.Binomial())
```

```
In [360]: result = full_logit_model.fit()
```

In [361]: activity_data

Out[361]:

	active	age	gender_M	metro_area_Birmingham	metro_area_Charlotte	metro_area_Detroit
0	1	30	0	0	0	1
1	1	43	0	0	1	0
2	1	41	0	0	0	0
3	1	56	0	0	0	0
4	1	33	0	0	0	1
...
5415	0	21	1	0	0	0
5416	0	42	0	1	0	0
5417	0	40	0	0	0	0
5418	0	40	0	0	0	0
5419	0	30	0	0	0	0

5420 rows × 13 columns

Apply this model to the same data that the model was trained on and assess the prediction accuracy.

In [362]: *# Applying the model on activity_data to assess the prediction accuracy*
 predictions = result.predict(activity_data[explanatory_cols])
 predictions

Out[362]: 0 0.561944
 1 0.469301
 2 0.586239
 3 1.000000
 4 0.505904
 ...
 5415 0.298635
 5416 0.542475
 5417 0.437380
 5418 0.437380
 5419 0.549623
 Length: 5420, dtype: float64

In [363]: *# Converting probability to binary category active = 1 if x >=0.5, active = 0 if x < 0.5*
 predictions_nominal = [1 if x > 0.5 else 0 for x in predictions]

In [364]: *# Calculating the accuracy of the model*
 from sklearn import metrics
 print('Accuracy: ', metrics.accuracy_score(activity_data['active'], predictions_nominal))

Accuracy: 0.5806273062730627

The accuracy of the model is 58%

5. Split the data into training and test samples, and build a model over the training data using the following Python code:

```
In [365]: # Splitting the data into training and test set
training_data = activity_data[1:4000]
test_data = activity_data[4001:].copy()

training_logit_model = sm.GLM(training_data['active'], training_data[explanatory_cols], family=sm.families.Binomial())

training_result = training_logit_model.fit()
```

```
In [366]: training_result.summary()
```

Out[366]: Generalized Linear Model Regression Results

Dep. Variable:	active	No. Observations:	3999			
Model:	GLM	Df Residuals:	3987			
Model Family:	Binomial	Df Model:	11			
Link Function:	logit	Scale:	1.0000			
Method:	IRLS	Log-Likelihood:	-2554.6			
Date:	Wed, 31 Mar 2021	Deviance:	5109.3			
Time:	22:48:23	Pearson chi2:	3.97e+03			
No. Iterations:	22					
Covariance Type:	nonrobust					
	coef	std err	z	P> z 	[0.025	0.975]
age	0.0071	0.003	2.227	0.026	0.001	0.013
gender_M	-0.5802	0.097	-5.981	0.000	-0.770	-0.390
metro_area_Birmingham	-0.1187	0.115	-1.037	0.300	-0.343	0.106
metro_area_Charlotte	-1.7834	0.382	-4.663	0.000	-2.533	-1.034
metro_area_Detroit	-0.1390	0.138	-1.007	0.314	-0.409	0.132
metro_area_Houston	-0.4865	0.112	-4.358	0.000	-0.705	-0.268
metro_area_Mobile	-1.7606	0.284	-6.202	0.000	-2.317	-1.204
metro_area_Nashville	21.8608	1.33e+04	0.002	0.999	-2.6e+04	2.6e+04
metro_area_Tampa	0.1892	0.127	1.484	0.138	-0.061	0.439
device_Mobile	-1.5818	0.291	-5.428	0.000	-2.153	-1.011
device_Tablet	-1.2830	0.298	-4.309	0.000	-1.867	-0.699
const	2.0245	0.327	6.187	0.000	1.383	2.666

```
In [367]: # Applying the model on test_data to assess the prediction accuracy
predictions_test = training_result.predict(test_data[explanatory_cols])
```

```
In [368]: # Converting probability to binary category active = 1 if x >=0.5, active = 0 if x < 0.5
predictions_test_nominal = [1 if x > 0.5 else 0 for x in predictions_test]
```

```
In [369]: # Calculating and printing the accuracy
print('Accuracy: ', metrics.accuracy_score(test_data['active'], predictions_test_nominal))
```

Accuracy: 0.21071176885130374

The accuracy of the model was 21%, which was less than the accuracy in ques.4. This is because in the previous question, the model was evaluated on the same data on which the model was trained i.e. the activity_data. So when we evaluate the model that we trained we get high scores, this shows how well our model learnt from our training data.

However, in question 5, the model is evaluated on a new data set and so the accuracy of the model is reduced.

One of the major reason for low accuracy is overfitting. Overfitting models the training data very well.

It takes place when a model learns the detail and noise in the training data well and negatively impacts the performance of the model on new data. This means that the noise or random fluctuations in the training data is picked up and learned as concepts by the model. But, these concepts do not apply to new data.

6.This data comes from a subset of userdata JSON blobs stored in our database. Parse out the values (stored in the “data_to_parse” column) into four separate columns. So for example, the four additional columns for the first entry would have values of “N”, “U”, “A7”, and “W”. You can use any R functions/packages you want for this.

```
In [370]: # Reading the csv file
InterviewData_Parsing = pd.read_csv('/Users/abhilashasinha/Downloads/interview_challenge_data_analyst/InterviewData_Parsing.csv')
```

```
In [371]: # Data in InterviewData_Parsing
InterviewData_Parsing
```

Out[371]:

	userid	data_to_parse
0	54f3ad9a29ada	"value":"N;U;A7;W"}}
1	54f69f2de6aec	"value":"N;U;l6;W"}}
2	54f650f004474	"value":"Y;U;A7;W"}}
3	54f52e8872227	"value":"N;U;l1;W"}}
4	54f64d3075b72	"value":"Y;U;A7;W"}}
...
948	54f5eb32d1a5b	"value":"N;U;A1;W"}}
949	54f34bd1a812a	"value":"N;C;A2;L"}}
950	54f34aa1e1f00	"value":"Y;U;A1;W"}}
951	54f47d97846bc	"value":"N;U;l4;L"}}
952	54f5337f14bd6	"value":"N;U;A4;W"}}

953 rows × 2 columns


```
In [372]: # Importing library 're' and using it to remove the special characters from the column data_to_parse
import re

InterviewData_Parsing.data_to_parse = InterviewData_Parsing.data_to_parse.apply(lambda x: ' '.join(re.findall(r'\w+', x)))
```

```
In [373]: # Splitting the column data_to_parse into the 4 additional columns
InterviewData_Parsing[['data_to_parse_val', 'col1', 'col2', 'col3', 'col4']] = pd.DataFrame([x.split(' ') for x in InterviewData_Parsing['data_to_parse'].tolist()])
```

```
In [374]: # Displaying the data after parsing
InterviewData_Parsing
```

Out[374]:

	userid	data_to_parse	data_to_parse_val	col1	col2	col3	col4
0	54f3ad9a29ada	value N U A7 W	value	N	U	A7	W
1	54f69f2de6aec	value N U I6 W	value	N	U	I6	W
2	54f650f004474	value Y U A7 W	value	Y	U	A7	W
3	54f52e8872227	value N U I1 W	value	N	U	I1	W
4	54f64d3075b72	value Y U A7 W	value	Y	U	A7	W
...
948	54f5eb32d1a5b	value N U A1 W	value	N	U	A1	W
949	54f34bd1a812a	value N C A2 L	value	N	C	A2	L
950	54f34aa1e1f00	value Y U A1 W	value	Y	U	A1	W
951	54f47d97846bc	value N U I4 L	value	N	U	I4	L
952	54f5337f14bd6	value N U A4 W	value	N	U	A4	W

953 rows × 7 columns

Additional Questions – Pick One

A) Within our web and mobile apps, members can generally find items through search and/or the product category tree (note that you can also search after clicking into a product category, in which case the search is filtered by the chosen category). Let's say that we decide to test a different product category tree. The Product team asks for your help in setting up the test and calling the results. How would you help them: (i) figure out how long we should run this test; (ii) decide what metric to measure; (iii) and then evaluate the test?

How would you help them: The new product category tree can be tested with the help of A/B testing. A/B testing refers to a randomized experimentation process where two or more versions of a web page or page elements are shown to different segments of website visitors at the same time to determine which version leaves the maximum impact and drives business metrics.

This is one of the simplest ways to understand the performance of any website using statistical analysis while spending less time and money.

i. Figure out how long we should run this test: The test should run at least for one complete week. This is because for a few websites the conversion rates can be low during weekdays and can increase over the weekends and vice-versa. Considering that the web and mobile app is for Shipt, it can be possible that working people visit the website during weekends, or stay-at-home mothers can visit the website during weekdays. So, to get valid test data, tests should run throughout the week so as to include all possible fluctuations. The duration will also depend on the website traffic. If the traffic is lower, the test will have to run for a longer time.

ii. Decide what metric to measure: One of the most important metrics will be click-through rate. This will give the percentage of people that clicked on the search product category. This helps to measure the success of marketing efforts. Other metrics would be the bounce rate i.e. the percentage of visitors who clicked on the product category but did not stay there and left, the conversion rate, and the number of people who added their products into the cart after searching.

iii. Evaluate the test?: In order to evaluate the test, the different metrics that have been measured through the Website Optimizer should be considered. For instance, a higher click-through rate shows that the search category is engaging and people are interested in clicking and navigating through the category. A 2% click-through-rate is usually considered good. Similarly, if the bounce rate is high then it shows that the visitors did not find the page or content attractive and so they did not stay for a long time.

In []: